

Location or Classification: Which better predicts commercial unit survival in Edinburgh, and do these factors interact?

B272011

27th March 2026

1 Overview

This report investigates whether geographic location or commercial classification better predicts the long-term survival of commercial units in Edinburgh, and whether these two factors interact. Understanding these patterns is valuable for urban planning, economic forecasting, and aspiring business owners.

The analysis draws on five shop surveys conducted by the City of Edinburgh Council between 1986 and 2015. The raw CSV files were merged on a unique unit identifier, a binary survival target variable was engineered (active in both 1986 and 2015), and a Region variable was extracted from postcodes. After filtering to units active in 1986 with a known postal region and removing sparse categories (regions and classes with ≤ 20 units), the working sample consisted of 6,712 units, of which 86.6% survived to 2015. Exploratory data analysis with dot maps and plots examined survival variation across both predictors. Postal district survival rates ranged from 65.9% (EH5) to 96.6% (EH17), a 30-percentage-point spread, whereas classification showed only a 10-point range. Chi-squared tests formally validated these associations. A logistic regression model was then fitted with Region and Class as one-hot encoded categorical predictors, and nested models were compared using the Likelihood Ratio Test (LRT) to determine relative predictor importance.

The LRT showed that geographic location is the substantially stronger predictor of survival compared to commercial classification. However, the model's modest predictive power suggests that unmeasured variables, such as rent and area-level deprivation, also play substantial roles in unit longevity. The interaction between Region and Class was not significant ($p = 0.093$), indicating the two factors operate independently. For aspiring business owners, choosing where to locate may matter more than choosing what to sell.

2 Introduction

Context and motivation This study focuses on the long-term dynamics of commercial unit survival. In the UK alone, 317,000 new businesses were established and 280,000 were shut down in 2024 [1], demonstrating the highly volatile commercial market. Between 1986 and 2015, the City of Edinburgh Council carried out five thorough shop surveys in order to obtain a long-term perspective on unit change within the city. Examining this dataset will show us the underlying categorical and geographic elements that determine a unit's longevity. Understanding these survival patterns offers valuable, data-driven insights for modern urban planning, local economic forecasting, and aspiring business owners navigating an increasingly competitive market.

Previous work Numerous studies have been conducted on the evolution and survival of retail businesses. A 2011 technical report by the City of Edinburgh Council analysed local shopping centres between 1986 and 2010, noting a 6% decrease in shop units and a noticeable shift towards non-retail services, with survival outcomes varying significantly between city centre and out-of-town commercial zones [2]. Expanding beyond Edinburgh, recent spatial studies demonstrated that a combination of internal business

characteristics (like store size and sales volume) and external geographic factors (such as local store density and neighbourhood demographics) significantly impact store longevity [3, 4]. Furthermore, recent media analysis shows that major Edinburgh retail hubs, like Princes Street, continue to experience rapid turnover based on shop category, with traditional retail declining while service-based shops expand [5]. This project builds upon these concepts by applying statistical modeling and data analysis to determine the exact extent of these geographic and categorical influences in Edinburgh.

Objectives This project sets out to answer the primary question: To what extent do factors such as commercial classification and geographic location influence the long-term survival rates of Edinburgh commercial units across the 1986–2015 period?

To address this, the analysis will pursue the following sub-objectives:

- Determine exactly how many commercial units successfully survive over this multi-decade period.
- Evaluate whether a unit’s geographic location or commercial classification matters for its survival over time.
- Quantify the specific impact of these factors using a logistic regression model to determine which commercial classifications or regional locations are most significant for long-term survival.
- Compare the relative importance of geographic location versus commercial classification to identify which feature serves as a stronger predictor of a unit’s lifespan.

3 Data

Data provenance The dataset used in this analysis consists of five historical shop surveys conducted and compiled by the City of Edinburgh Council in 1986, 1996, 2004, 2010, and 2015. The data was obtained by downloading the CSV files directly from the Open Data Scotland website ¹. Because these datasets are published as open government data under the Open Government Licence, they are publicly available for reuse, which fully permits their modification and analysis for academic research [6].

Data description The raw data is distributed across five separate CSV files corresponding to the survey years. The number of recorded commercial units fluctuates across the decades, with a total of 7,668 unique commercial units. Each file contains between 10 and 11 variables capturing spatial coordinates (e.g., X, Y), business names and descriptions, commercial use classifications (e.g., USECLASS), address details, and a unique OBJECTID. However, the raw data presents significant structural challenges. Certain fields are messy or incomplete, and the exact definition of legacy features requires careful interpretation. Most notably, the critical “region” variable required for spatial analysis is not natively present, meaning it must be externally engineered using the provided postcodes.

The USECLASS variable follows the statutory Scottish town planning classification system, which we have categorized into distinct business types. This includes Class 1 (Shops), Class 2 (Financial and Professional Services), Class 3 (Food and Drink), Class 4 (Business), and Class 11 (Assembly and Leisure), alongside *Sui Generis* units which are uniquely classified [7]. Although Class 1 and 2 were merged into a new Class 1A in March 2023, they are retained separately here to accurately reflect the historical survey data [7].

Data processing The five datasets were combined into a continuous record for every commercial unit using the distinct OBJECTID. Missing values were sequentially filled using the most recent non-null entry across survey years to maximise data retention. Text fields were standardised to uppercase, and empty

¹<https://opendata.scot>

strings were converted to NaN values. A new Region variable was created by extracting the postal district from the POSTCODE column (for example, “EH8”). Lastly, a binary ‘Survived’ target variable was engineered. A unit that was operational in 1986 was given a “1” if it continued commercial operations in 2015, and a “0” if recorded as missing, vacant, or demolished. The final dataset did not include rows that lacked crucial regional data. Finally, a sparsity filter was applied to exclude commercial classes or regions with fewer than 20 units; this step was necessary to prevent algorithmic instability and unreliable estimates during the logistic regression.

4 Exploration and analysis

After filtering to units active in 1986 with a known postal region, and excluding sparse categories (≤ 20 units), the working sample consisted of 6,712 units. Of these, 5,810 (86.6%) survived to 2015 and 902 (13.4%) did not. Before conducting statistical modelling, we evaluate our two key predictors of survival: postal district (Region) and initial commercial classification (Class).

Geographic analysis (Figure 1) shows a stark, 30-percentage-point difference in long-term survival between areas. Suburban residential districts like EH17 (96.6%) overwhelmingly outperformed highly-squeezed, tourism-heavy inner-city areas like EH8 (74.6%) [8] and regenerating industrial zones like EH5 (65.9%) [9]. A plausible hypothesis for this high retention is that commercial units in these neighbourhoods primarily serve stable local populations with long-term, everyday needs.




Figure 1: Survival rate of commercial units across Edinburgh’s postal regions (1986–2015). Dot colour encodes survival rate (purple–orange–yellow) and dot size is proportional to the number of units in that region.

Similarly, Figure 2 confirms survival varies by initial use-class. Units offering essential provisions or hospitality (e.g., Food and Drink (Class 3) at 90.7%; Shops (Class 1) at 87.0%) proved far more resilient than discretionary or specialised sectors like Assembly and Leisure (Class 11 at 80.4%). This is likely

because food and shop units serve essential daily needs, whereas leisure spaces with higher overheads and reliance on disposable income are more vulnerable to economic downturns.




Figure 2: Survival rate of commercial units by their initial 1986 use-class classification. Each label shows the number of units in that category. The dashed line marks the overall mean survival rate (86.6%).

To formally validate these observed associations, we applied Chi-squared tests of independence (Table 1). Both Region and Class yielded significant results ($p < 0.05$), statistically confirming their validity as distinct predictors for our subsequent logistic regression modelling.

Table 1: Chi-Squared Test of Independence Results

Predictor Variable	Chi-Squared (χ^2)	d.f.	p-value
Postal Region	217.03	17	1.05×10^{-36}
1986 Commercial Class	14.35	5	1.35×10^{-2}

To quantify these effects and address our core research question, we fitted a logistic regression model predicting the binary survival outcome from Region and Class as categorical predictors. These features were one-hot encoded, with **Class 1 (Shop)** and **EH1 (City Centre)** withheld as the reference categories to avoid the dummy-variable trap. The model was trained on an 80/20 stratified split using both `statsmodels` (for statistical inference) and `sklearn` (with inverse class-frequency weighting to account for the 86.6% majority-class imbalance).

Figure 3 displays the estimated coefficients with 95% confidence intervals. Regions whose intervals exclude zero are statistically significant predictors. EH5 – North (Coastal: Granton, Trinity, Pilton) and EH8 – Inner East (Southside and Northfield) show the strongest negative effects on survival odds, while EH10 and EH7 (suburban areas) show significant positive effects. Among commercial classes, Sui Generis units have significantly lower survival odds, whereas CLASS 3 (Food and Drink) shows a positive, but not significant trend.




Figure 3: Logistic regression coefficients (log-odds) with 95% confidence intervals. Predictors whose intervals exclude the dashed zero line are statistically significant at the 5% level.

To determine whether Region or Class is the stronger predictor, we compared nested models using the Likelihood Ratio Test (Table 2). This test quantifies whether adding or dropping a specific variable significantly changes the model’s fit to data by comparing our full complex model against simpler, baseline versions. Dropping Region from the full model produces a far larger LRT statistic (150.82, $p = 1.69 \times 10^{-23}$) than dropping Class (11.61, $p = 0.0405$), confirming that **geographic location is the substantially stronger predictor** of long-term survival. Adding Region $\times$ Class interaction terms did not significantly improve model fit (LRT = 102.69, $p = 0.093$), indicating that an additive model is sufficient and the two predictors operate largely independently. In practical terms, the contextual factors captured by location (footfall, demographic stability, and regeneration pressure) matter more for a unit’s longevity than what type of business/class initially occupies it. For aspiring business owners, choosing *where* to open may therefore be a more consequential decision than choosing *what* to open.

Table 2: Likelihood Ratio Test comparison of nested logistic regression models, each compared against the full additive model (Region + Class) as the reference.

Model	LRT	p-value
Region-only (drop Class)	11.61	0.0405
Class-only (drop Region)	150.82	1.69×10^{-23}
Interaction (Region $\times$ Class)	102.69	0.093

5 Discussion and conclusions

Summary of findings This analysis addressed whether geographic location or commercial classification better predicts long-term commercial unit survival in Edinburgh between 1986 and 2015. Of 6,712 units

active in 1986, the vast majority (86.6%) survived to 2015. Both Region and Class were confirmed as statistically significant predictors via chi-squared tests, and logistic regression quantified their relative importance. The Likelihood Ratio Test showed that Region is the substantially stronger predictor, while the interaction between the two was not significant, meaning the effects of location and commercial classification operate independently.

Evaluation of own work: strengths and limitations A key strength of this work is the use of formal model comparison via the LRT rather than relying solely on descriptive statistics, which allows us to make rigorous claims about relative predictor importance. The use of inverse class-frequency weighting also ensured the model did not simply default to predicting the majority class.

However, several limitations affect the conclusions. The model achieved a modest test ROC-AUC of 0.655 (Figure 4), indicating that Region and Class alone explain only a small share of the variation in survival. Important unmeasured variables such as rent, proximity to transport, area-level deprivation, and individual management decisions likely play substantial roles. Additionally, the survival variable tracks whether a physical *unit* remained commercially active, not whether the same *business* persisted. A unit where the original shop closed and a new one opened would still be coded as “survived,” potentially overstating true business longevity. Finally, some postal districts and use-classes contained very few observations, leading to wide confidence intervals on those coefficients.

Comparison with any other related work

Our finding that location outweighs commercial classification aligns with the City of Edinburgh Council’s 2010 report, which noted significant variation in survival between city-centre and out-of-town zones [2]. It is also consistent with Lee et al. [4], who found that geographic context, including local store density and neighbourhood characteristics was a primary driver of small business survival. Our contribution is to quantify this relationship specifically for Edinburgh over a longer time horizon and to formally test the interaction, finding that the two factors do not combine in unexpected ways.

Improvements and extensions

With more time, three improvements could strengthen this analysis. First, incorporating the intermediate survey waves (1996, 2004, 2010) as longitudinal features, such as whether a unit was active at each wave or how many times it changed use-class would give the model a richer picture of each unit’s trajectory rather than relying solely on its 1986 starting state. Second, a non-linear model such as a Random Forest could capture complex interactions without requiring explicit interaction terms. Third, linking the dataset to external sources such as the Scottish Index of Multiple Deprivation would allow the inclusion of socioeconomic variables like local income levels and population density, which likely account for much of the unexplained variance.




Figure 4: ROC curves for the additive and interaction logistic regression models on the held-out test set. Both models perform above chance ($AUC > 0.5$) but show limited discriminative power with only Region and Class as predictors.

References

- [1] House of Commons Library. *Business Statistics*. Accessed: 2026-02-20. 2024. URL: <https://commonslibrary.parliament.uk/research-briefings/sn06152/>.
- [2] City of Edinburgh Council. *Edinburgh's Shopping Centres 1986–2010*. Accessed: 2026-02-20. 2010. URL: https://www.broughtonspurtle.org.uk/sites/default/files/story/edinburghs_shopping_centres.pdf.
- [3] Joshua L. Warren and Penny Gordon-Larsen. 'Factors associated with supermarket and convenience store closure: a discrete time spatial survival modelling approach'. In: (2018). Accessed: 2026-02-20. URL: <https://pmc.ncbi.nlm.nih.gov/articles/PMC5990034/>.
- [4] DongHyeon Lee et al. 'Classification of commercial districts based on predicting the survival rate of food service market in Seoul'. In: (2025). Accessed: 2026-02-20. URL: <https://pmc.ncbi.nlm.nih.gov/articles/PMC12244476/>.
- [5] BBC News. *How lockdown changed Scotland's high streets*. Accessed: 2026-02-25. 2022. URL: <https://www.bbc.co.uk/news/uk-scotland-63865960>.
- [6] The National Archives. *Open Government Licence for Public Sector Information, Version 3.0*. Accessed: 2026-03-18. 2014. URL: <https://www.nationalarchives.gov.uk/doc/open-government-licence/version/3/>.
- [7] City of Edinburgh Council. *Quick Guide to Use Classes and Permitted Changes of Use*. Accessed: 2026-03-12. 2024. URL: <https://www.edinburgh.gov.uk/downloads/file/34675/quick-guide-to-changes-of-use>.
- [8] Wikipedia. *EH postcode area*. Accessed: 2026-03-14. 2025. URL: https://en.wikipedia.org/wiki/EH_postcode_area.
- [9] Wikipedia. *Granton Waterfront Regeneration*. Accessed: 2026-03-12. 2025. URL: https://en.wikipedia.org/wiki/Granton_Waterfront_regeneration.

Predicting Edinburgh Commercial Unit Survival (1986–2015) Using Logistic Regression

Imports

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from functools import reduce
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.metrics import (accuracy_score, classification_report, roc_curve, auc)
from scipy.stats import chi2_contingency
import warnings
from matplotlib.lines import Line2D
from scipy.stats import chi2
import matplotlib.cm as cm
import matplotlib.colors as mcolors
import matplotlib.image as mpimg

# Global plotting: whitegrid theme, colourblind-friendly palette, legible font
sns.set_theme(style="whitegrid")
sns.set_palette("colorblind")
plt.rcParams.update({'font.size': 14, 'figure.dpi': 200})
```

Data

Loading the Survey Data

We load each file and set `OBJECTID` as the index so we can align records across survey years when constructing the master dataset.

```
In [2]: df_86 = pd.read_csv('datasets/Shop_Survey_1986.csv').set_index('OBJECTID')
df_96 = pd.read_csv('datasets/Shop_Survey_1996.csv').set_index('OBJECTID')
df_04 = pd.read_csv('datasets/Shop_Survey_2004.csv').set_index('OBJECTID')
df_10 = pd.read_csv('datasets/Shop_Survey_2010.csv').set_index('OBJECTID')
df_15 = pd.read_csv('datasets/Shop_Survey_2015.csv').set_index('OBJECTID')

df_86.head().T
```

Out [2]:	OBJECTID	1	2	3	4	5
	X	326794.4614	327067.4706	326939.8169	328372.6636	325432.0
	Y	673934.3841	674402.0065	674274.0781	674066.0148	674403.0
	DESCRIP86	restaurant	hairdresser	vacant	surgery	restaurant
	BUSINESS86					
	USECLASS86	Class 3	Class 1	Vacant	Class 2	Class 3
	USECLASS96	Class 9	Class 1	Demolished	Class 4	Class 3
	USECLASS04	Class 3			Class 4	
	USECLASS10	Class 3			Class 4	
	ADDRESS	9 ABBEY STRAND		4 ABERCORN ROAD		
	POSTCODE	EH8 8DU		EH8 7DE		

The 2015 survey uses `CVR15_POST` instead of `POSTCODE` . We rename this column to harmonise naming conventions before merging, ensuring the postcode field can be coalesced consistently across all five survey years.

```
In [3]: df_15 = df_15.rename(columns={'CVR15_POST': 'POSTCODE'})
df_15['POSTCODE'].head()
```

```
Out[3]: OBJECTID
1      EH8 8HH
2
3      EH8 8DU
4      EH8 8EA
5      EH8 7DE
Name: POSTCODE, dtype: str
```

Handling Whitespace-Only Values

Some raw CSV cells contain only whitespace (e.g. " "), which Pandas reads as non-null strings. These must be converted to `NaN` so that our subsequent `combine_first` merge correctly prioritises genuine non-empty values from other survey years when filling missing data.

```
In [4]: df_86, df_96, df_04, df_10, df_15 = [
        df.replace(r'^\s*$', np.nan, regex=True)
        for df in [df_86, df_96, df_04, df_10, df_15]
    ]

df_86.head().T
```

Out [4]:

OBJECTID	1	2	3	4	5
X	326794.4614	327067.4706	326939.8169	328372.6636	325432.0
Y	673934.3841	674402.0065	674274.0781	674066.0148	674403.0
DESCRIP86	restaurant	hairdresser	vacant	surgey	restaurant
BUSINESS86	NaN	NaN	NaN	NaN	NaN
USECLASS86	Class 3	Class 1	Vacant	Class 2	Class 3
USECLASS96	Class 9	Class 1	Demolished	Class 4	Class 3
USECLASS04	Class 3	NaN	NaN	Class 4	NaN
USECLASS10	Class 3	NaN	NaN	Class 4	NaN
ADDRESS	9 ABBEY STRAND	NaN	NaN	4 ABERCORN ROAD	NaN
POSTCODE	EH8 8DU	NaN	NaN	EH8 7DE	NaN

Merging Into a Master Dataset

We merge the five survey DataFrames into a single master dataset with one row per commercial unit. Shared columns (X , Y , POSTCODE , ADDRESS) are coalesced using `combine_first` , cascading through survey years to take the first non-null value and maximise data retention. Year-specific use-class and business name columns are joined to produce a wide-format table capturing each unit's trajectory across all survey waves.

```
In [5]: survey_dfs = [df_86, df_96, df_04, df_10, df_15]
shared_cols = ['X', 'Y', 'POSTCODE', 'ADDRESS']

df_master = pd.DataFrame(index=reduce(pd.Index.union,
                                     [df.index for df in survey_dfs]))

df_master[shared_cols] = reduce(
    lambda a, b: a.combine_first(b),
    [df[shared_cols] for df in survey_dfs]
)

df_master = df_master.assign(
    Class_1986=reduce(lambda a, b: a.combine_first(b),
                     [df['USECLASS86'] for df in [df_86, df_96, df_04, df_10]]),
    Class_1996=reduce(lambda a, b: a.combine_first(b),
                     [df['USECLASS96'] for df in [df_96, df_86, df_04, df_10]]),
    Class_2004=reduce(lambda a, b: a.combine_first(b),
                     [df['USECLASS04'] for df in [df_04, df_86, df_96, df_10]]),
    Class_2010=reduce(lambda a, b: a.combine_first(b),
                     [df['USECLASS10'] for df in [df_10, df_86, df_96, df_04]]),
    Class_2015=df_15['USECLASS15'],
)
```

```
In [6]: df_master = df_master.join(
    pd.concat([
        df_86['BUSINESS86'].rename('Business_1986'),
        df_96['BUSINESS96'].rename('Business_1996'),
        df_04['BUSINESS04'].rename('Business_2004'),
        df_10['BUSINESS10'].rename('Business_2010'),
        df_15['BUSINESS15'].rename('Business_2015'),
    ], axis=1)
)
```

```
df_master = df_master.rename_axis('OBJECTID').reset_index()
```

```
In [7]: print(f"Master dataset shape: {df_master.shape}")
display(df_master.head().T)
```

Master dataset shape: (7668, 15)

	0	1	2	3	4
OBJECTID	1	2	3	4	5
X	326794.4614	327067.4706	326939.8169	328372.6636	325432.0
Y	673934.3841	674402.0065	674274.0781	674066.0148	674403.0
POSTCODE	EH8 8DU	EH8 8DU	EH8 8EA	EH8 7DE	EH3 6LB
ADDRESS	9 ABBEY STRAND	9(GF) ABBEY STRAND	55 ABBEYHILL	4 ABERCORN ROAD	8B ABERCROMBY PLACE
Class_1986	Class 3	Class 1	Vacant	Class 2	Class 3
Class_1996	Class 9	Class 1	Demolished	Class 4	Class 3
Class_2004	Class 3	Class 1	UNKNOWN USE	Class 4	Class 3
Class_2010	Class 3	Class 1	Vacant	Class 4	Class 3
Class_2015	Class 1	Class 3	Class 1	Vacant	Class 2
Business_1986	NaN	NaN	NaN	NaN	NaN
Business_1996	NaN	L & G Hair	NaN	Foundations	Premier Mandarin Cuisine
Business_2004	Café at the Palace	Queen's Gallery Shop	NaN	Foundations	Bellini
Business_2010	Café at the Palace	Queen's Gallery Shop	vacant	Ardmore Home Care	Karens Unicorn
Business_2015	BONK & CO LTD	Caf, at the Palace	Queen's Gallery Shop	vacant	Click Printwear

Data Cleaning & Feature Engineering

Text fields are stripped of whitespace and converted to uppercase to prevent spurious duplicates (e.g. "a1" vs "A1"). Crucially, the **Region** variable is not natively present in the data and must be engineered by extracting the postal district from the **POSTCODE** column.

```
In [8]: text_cols = df_master.select_dtypes(include=['object', 'str']).columns
df_master[text_cols] = (df_master[text_cols]
                        .apply(lambda s: s.str.strip().str.upper()))

df_master['Region'] = df_master['POSTCODE'].str.split(' ').str[0]
df_master[['Region']].head()
```

Out [8]:

	Region
0	EH8
1	EH8
2	EH8
3	EH8
4	EH3

Defining Target Variable: Unit Survival

To operationalise our research question, we identify units that were commercially active in 1986 (excluding those classified as `VACANT` , `DEMOLISHED` , or `UNKNOWN USE`) and encode a binary `Survived` target: 1 if the unit remained active in 2015, 0 otherwise. Units lacking a postal region are dropped since Region is a key predictor. After this filtering, the working sample consists of **6,742 units**.

```
In [9]: inactive_labels = ['VACANT', 'DEMOLISHED', 'UNKNOWN USE']

active_86 = (df_master.query("Class_1986.notna() and Class_1986 "
                             "not in @inactive_labels").copy())

active_86['Survived'] = (
    active_86['Class_2015'].notna() &
    ~active_86['Class_2015'].isin(inactive_labels)
).astype(int)

df_analysis = active_86.dropna(subset=['Region']).copy()
```

```
In [10]: print(f"Units active in 1986 with known region: {len(df_analysis):,}")
df_analysis[['OBJECTID', 'Survived']].head()
```

Units active in 1986 with known region: 6,742

Out [10]:

	OBJECTID	Survived
0	1	1
1	2	1
3	4	0
4	5	1
10	11	1

Data Refinement: Handling Sparse Categories

Before proceeding to our final statistical tests and logistic regression modeling, we must address "sparse categories" (regions or commercial classes that contain very few observations).

In logistic regression, categories with extremely low counts (e.g., a region with only one or two units) can lead to perfect separation. This occurs when a predictor perfectly predicts the outcome by chance, causing the model's coefficients to explode toward infinity and preventing the Maximum Likelihood Estimation (MLE) solver from converging. To ensure algorithmic stability and the reliability of our coefficients, we implement a threshold: only regions and commercial classes with at least 20 observations are retained in our final analytical sample.

```
In [11]: region_counts = df_analysis['Region'].value_counts()
class_counts = df_analysis['Class_1986'].value_counts()

valid_regions = region_counts[region_counts >= 20].index
valid_classes = class_counts[class_counts >= 20].index
```

```
In [12]: df_analysis = df_analysis[
    (df_analysis['Region'].isin(valid_regions)) &
    (df_analysis['Class_1986'].isin(valid_classes))
].copy()

total_units = len(df_analysis)
survived_counts = df_analysis['Survived'].value_counts()
survived_pct = (survived_counts.get(1, 0) / total_units) * 100
did_not_survive_pct = 100 - survived_pct

print(f"New working sample size: {total_units},")
print(f"Survived: {survived_counts.get(1, 0)} ({survived_pct:.1f}%)")
print(f"Did not survive: {survived_counts.get(0, 0)} ({did_not_survive_pct:.1f}%)")
```

New working sample size: 6,712
Survived: 5810 (86.6%)
Did not survive: 902 (13.4%)

Interpretation of Filtered Data

By applying this filter, we have refined our working sample to 6,712 units, maintaining a baseline survival rate of 86.6%.

This cleaned dataset will be used for all subsequent analyses, including our Chi-Squared tests and Logistic Regression models. Removing these unstable edge cases allows us to generate narrower confidence intervals and more trustworthy p-values, ensuring that the results discussed in the final report reflect meaningful trends rather than the noise of individual units.

Exploratory Data Analysis

With the master dataset constructed and the target variable defined, we can dive deeper into the data. A first look at the raw records reveals a nuance worth understanding before any analysis. The same address can appear multiple times in the table.

```
In [13]: print(f"Master dataset shape: {df_master.shape}")
display(df_master[['X', 'Y', 'ADDRESS']].head(15))
```

Master dataset shape: (7668, 16)

	X	Y	ADDRESS
0	326794.4614	673934.3841	9 ABBEY STRAND
1	327067.4706	674402.0065	9(GF) ABBEY STRAND
2	326939.8169	674274.0781	55 ABBEYHILL
3	328372.6636	674066.0148	4 ABERCORN ROAD
4	325432.0000	674403.0000	8B ABERCROMBY PLACE
5	327065.6429	675875.6072	4 ACADEMY STREE
6	327068.7502	675878.9646	1 ALBANY STREET
7	326540.7538	676593.4503	13(GF) ALBANY STREET
8	326545.0261	676599.0733	2 ALBANY STREET
9	326547.9528	676602.6754	1 ALBANY STREET
10	325604.9710	674390.9144	1 ALBANY STREET
11	325596.7830	674420.3208	2 ALBANY STREET
12	325600.3632	674421.5142	4 ALBANY STREET
13	326504.1046	675014.6623	1 ALBERT PLACE
14	326499.3377	675008.6754	3 ALBERT PLACE

Why the Same Address Appears Multiple Times

Each row represents a distinct **commercial unit** (an individual lettable premises), not a building. Duplicate addresses arise for two distinct reasons:

1. **Multiple units within one building.** A single building can contain several separately-tenanted units (a ground-floor shop, a basement café, an upper-floor office) each assigned its own `OBJECTID` and tracked independently across all five survey years. When a floor or sub-unit qualifier is recorded it is embedded directly in the address string (e.g. `9(GF) ABBEY STRAND` for the ground-floor unit vs. `9 ABBEY STRAND` for another unit in the same building); when no qualifier is present the units share an identical address string.
2. **Data quality issues.** Some apparent duplicates are artifacts of incomplete or inconsistent data entry. For example, `1 ALBANY STREET` appears three times with coordinates up to hundreds of metres apart (far too distant to represent units in the same building) suggesting that unrelated units had their address field recorded incorrectly or left incomplete. Such cases should be treated with caution in any address-level analysis.

Survival Count

We can also have a look at how many units survived throughout the multi-decade period from 1986 to 2015.

```
In [14]: df_analysis['Survived'].value_counts()
```

```
Out[14]: Survived
1      5810
0       902
Name: count, dtype: int64
```

We see that **5,810 (86.6%) survived** to 2015 and **902 (13.4%) did not**. This establishes a majority-class baseline of 86.6% accuracy that our predictive model must exceed to prove informative. Furthermore, this highlights a significant class imbalance in our target variable, which we will need to account for during modelling (e.g., by applying balanced class weights) to ensure our model doesn't simply learn to always predict survival.

With the overall survival baseline established, we now investigate whether survival rates differ systematically across the two categorical predictors: postal region and initial commercial classification. Beginning with geographic location.

Survival Rate by Postal Region

We compute survival rates grouped by Edinburgh postal region. This addresses the first dimension of our research question: whether geographic location is associated with long-term commercial unit survival.

```
In [15]: # Group, Aggregate, and Format
region_summary = (
    df_analysis.query("Region in @valid_regions")
    .groupby('Region')['Survived']
    .agg(['mean', 'count'])
    .rename(columns={'mean': 'Survival_Rate'})
    .sort_values(by='Survival_Rate', ascending=False)
)

region_summary
```

Out [15]:

	Survival_Rate	count
Region		
EH17	0.966102	59
EH13	0.944444	54
EH10	0.941788	481
EH15	0.929752	242
EH7	0.916667	492
EH9	0.912752	298
EH16	0.907285	151
EH6	0.905054	653
EH2	0.880716	503
EH1	0.875740	676
EH4	0.873684	380
EH11	0.863063	555
EH3	0.854358	872
EH30	0.788462	52
EH12	0.788360	378
EH14	0.786070	201
EH8	0.746269	536
EH5	0.658915	129

```
In [16]: overall_rate = df_analysis['Survived'].mean() * 100
overall_rate
```

Out [16]: np.float64(86.56138259833135)

```
In [17]: postcode_coords = {
    'EH1': (55.9533, -3.2050),
    'EH2': (55.9550, -3.2200),
    'EH3': (55.9650, -3.2250),
    'EH4': (55.9725, -3.2900),
    'EH5': (55.9800, -3.2300),
    'EH6': (55.9800, -3.1900),
    'EH7': (55.9600, -3.1600),
    'EH8': (55.9480, -3.1750),
    'EH9': (55.9380, -3.2050),
    'EH10': (55.9300, -3.2350),
    'EH11': (55.9425, -3.2400),
    'EH12': (55.9450, -3.3100),
    'EH13': (55.9050, -3.2600),
    'EH14': (55.8950, -3.3250),
    'EH15': (55.9500, -3.1150),
    'EH16': (55.9280, -3.1550),
    'EH17': (55.9050, -3.1650),
    'EH30': (55.9900, -3.4000),
}

# Geographic extent of the PNG map [lon_min, lon_max, lat_min, lat_max]
map_extent = [-3.44, -3.04, 55.875, 56.005]
```

```

region_stats=(
    df_analysis.query("Region in @valid_regions")
    .groupby('Region')['Survived']
    .agg(['mean', 'count'])
)

rates = region_stats['mean'] * 100
sizes = region_stats['count'] * 2.0

lons = [postcode_coords[r][1] for r in region_stats.index]
lats = [postcode_coords[r][0] for r in region_stats.index]

norm = mcolors.Normalize(vmin=rates.min() - 2, vmax=rates.max() + 2)
cmap = cm.plasma

fig, ax = plt.subplots(figsize=(14, 8))


# Load and display the PNG as background
bg_img = mpimg.imread('edinburgh_map.png')
ax.imshow(bg_img, extent=map_extent, aspect='auto', zorder=0)

sc = ax.scatter(lons, lats,
                c=rates, cmap=cmap, norm=norm,
                s=sizes,
                edgecolors='black', linewidths=0.5,
                alpha=0.85, zorder=3)

for region, lon, lat in zip(region_stats.index, lons, lats):
    ax.annotate(region,
                xy=(lon, lat),
                fontsize=12, fontweight='bold',
                ha='center', va='bottom',
                xytext=(0, 8), textcoords='offset points')

cbar = fig.colorbar(sc, ax=ax, shrink=0.6, pad=0.02)
cbar.set_label('Survival Rate (%)', fontsize=16)
cbar.ax.tick_params(labelsize=13)
ax.set_axis_off()
ax.set_title(
    'Edinburgh Commercial Unit Survival Rate'
    ' by Postcode District (1986-2015)', fontsize=16, pad=10
)
fig.savefig('survival_by_region.png', dpi=200, bbox_inches='tight')
plt.show()

```



Interpretation: Survival rates span a **30-percentage-point gap** across regions. Suburban districts (e.g., EH17, EH13) consistently outperform inner-city areas, likely because they serve stable local populations with everyday needs. Conversely, areas experiencing major post-industrial regeneration (EH5) or intense tourism and rent pressures (EH8) exhibit the lowest survival rates.

Chi-Squared Test of Independence — Region

The dot map suggests that survival rates vary meaningfully by region, but visual inspection alone cannot distinguish a true effect from random sampling noise. To formally test this, we perform a **chi-squared test of independence**. This tests the null hypothesis that survival outcome and postal region are completely unrelated.

- H_0 (**Null Hypothesis**): Unit survival is independent of its postal region.
- H_1 (**Alternative Hypothesis**): There is a statistically significant relationship between postal region and unit survival.

Rejecting the null hypothesis at the 5% significance level provides formal statistical evidence that **Region** is a meaningful predictor, serving as a robust justification for including it as a feature in our upcoming logistic regression model.

```
In [18]: # Build a contingency table: rows = Region, columns = Survived (0 or 1)
contingency = pd.crosstab(
    df_analysis.query("Region in @valid_regions")['Region'],
    df_analysis.query("Region in @valid_regions")['Survived']
)

chi2_stat, p, dof, expected = chi2_contingency(contingency)
print(f"Chi-squared statistic : {chi2_stat:.2f}")
print(f"Degrees of freedom   : {dof}")
print(f"p-value                  : {p:.2e}")
```

Chi-squared statistic : 217.03
Degrees of freedom : 17
p-value : 1.05e-36

Interpretation: The test yields $\chi^2 = 217.03$ with 17 degrees of freedom (and $p = 1.05 \times 10^{-36}$), leading us to decisively reject the null hypothesis. This confirms that survival outcome and postal region are **not independent**. The geographic differences observed in the chart are statistically significant and not merely an artifact of sampling variability, formally justifying the inclusion of Region as a predictor in our logistic regression model.

Survival Rate by Initial Commercial Classification (1986–2015)

Next, we examine our second key predictor: a unit's initial 1986 commercial classification (Class). Under Scottish planning regulations, this variable categorises a unit's core commercial function, such as Class 1 (Shops), Class 3 (Food and Drink), or Sui Generis (unique/other units).

Because different commercial sectors possess varying degrees of resilience to long-term economic shifts, exploring this variable allows us to identify which types of businesses were most likely to survive the 29-year study period, and which were most vulnerable.

```
In [19]: class_summary = (  
    df_analysis  
    .groupby('Class_1986')['Survived']  
    .agg(['mean', 'count'])  
    .sort_values(by='mean', ascending=False)  
)  
  
class_summary
```

```
Out [19]:
```

	mean	count
Class_1986		
CLASS 3	0.906907	333
CLASS 1	0.870502	4363
CLASS 2	0.861301	584
CLASS 4	0.858195	543
SUI GENERIS	0.835113	843
CLASS 11	0.804348	46

```
In [20]: class_label_map = {  
    'CLASS 1': 'Class 1 – Shop',  
    'CLASS 2': 'Class 2 – Financial & Professional Services',  
    'CLASS 3': 'Class 3 – Food & Drink',  
    'CLASS 4': 'Class 4 – Business',  
    'CLASS 11': 'Class 11 – Assembly & Leisure',  
    'SUI GENERIS': 'Sui Generis – Unique/Other',  
}  
  
class_stats = (  
    df_analysis  
    .groupby('Class_1986')['Survived']  
    .agg(['mean', 'count'])  
    .sort_values('mean')  
)
```

```

cb = sns.color_palette("colorblind")
fig, ax = plt.subplots(figsize=(12, 7))

y_pos = range(len(class_stats))
labels = [f"{class_label_map.get(c, c)} ({int(n)} units)"
          for c, n in zip(class_stats.index, class_stats['count'])]
rates = class_stats['mean'] * 100


ax.scatter(rates, y_pos, color=cb[0], s=60, zorder=3)
ax.hlines(y_pos, xmin=0, xmax=rates, color=cb[0], linewidth=1.2, alpha=0.4)

ax.axvline(x=86.6, color=cb[3], linestyle='--', linewidth=1.5,
          label='Overall mean (86.6%)')

ax.set_yticks(y_pos)
ax.set_yticklabels(labels, fontsize=18)
ax.set_xlabel('Survival Rate (%)', fontsize=18)
ax.set_title('Commercial Unit Survival Rate by 1986 Classification (1986–2015)',
            fontsize=18, pad=10)
ax.legend(loc='lower right', fontsize=15)
ax.tick_params(axis='x', labelsize=18)
ax.set_xlim(60, 100)
ax.margins(y=0.05)

fig.savefig('survival_by_class.png', dpi=200, bbox_inches='tight')
plt.show()

```



Interpretation: Survival rates across commercial classes span a much narrower **10-percentage-point gap**. Essential services like Food & Drink (90.7%) and Shops (87.0%) prove highly resilient, whereas discretionary sectors like Assembly & Leisure (80.4%) are more vulnerable. Crucially, comparing this modest 10-point spread to the dramatic 30-point regional gap provides our first hint that where a business locates matters more than what it sells. A hypothesis we will formally test via the Likelihood Ratio Test.

Chi-Squared Test — Commercial Classification

We apply the same chi-squared test of independence to `Class_1986` to determine whether the observed variation in survival across use-classes is statistically significant or could have arisen by chance.

```

In [21]: # Contingency table: rows = Class_1986, columns = Survived (0 or 1)
contingency_class = pd.crosstab(df_analysis['Class_1986'], df_analysis['Survived'])

```

```
chi2_stat, p_class, dof_class, expected_class = chi2_contingency(contingency_class)
print(f"Chi-squared statistic : {chi2_stat:.2f}")
print(f"Degrees of freedom : {dof_class}")
print(f"p-value : {p_class:.2e}")
```

```
Chi-squared statistic : 14.35
Degrees of freedom : 5
p-value : 1.35e-02
```

Interpretation: The test yields $\chi^2 = 14.35$ with 5 degrees of freedom (and $p = 0.0135$), confirming that survival and commercial classification are **not independent**. Both Region and Class are now formally validated as significant predictors, justifying their joint inclusion in the logistic regression model.

Data Communication and Analysis

Logistic Regression

Why Logistic Regression?

Our response variable is binary (survived vs. did not survive), making logistic regression the natural modelling choice. It models the **log-odds** of survival as a linear combination of predictors, ensuring predicted probabilities remain in $[0, 1]$:

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x_1 + \dots + \beta_k x_k$$

We use two complementary implementations: `statsmodels.Logit` for statistical inference (coefficients, p-values, confidence intervals) and `sklearn.LogisticRegression` with `class_weight='balanced'` for fairer prediction that does not default to the 86.6% majority class. We fit both an **additive model** (Region + Class) and an **interaction model** (Region $\times$ Class) to test whether the effect of commercial classification depends on location.

Feature Preparation: One-Hot Encoding and Interaction Terms

Both predictors are categorical, so we convert them to binary dummy variables via one-hot encoding with `drop_first=True` to avoid the **dummy-variable trap**. Interaction terms (pairwise products of every Region dummy with every Class dummy) are also constructed to test whether certain business types fare differently in specific locations.

```
In [22]: # One-hot encode categoricals
X = pd.get_dummies(df_analysis[['Region', 'Class_1986']], drop_first=True, dtype=int)
y = df_analysis['Survived']

# Build interaction feature matrix (Region x Class)
region_cols_feat = [c for c in X.columns if c.startswith('Region_')]
class_cols_feat = [c for c in X.columns if c.startswith('Class_1986_')]

interaction_parts = []
for rc in region_cols_feat:
    for cc in class_cols_feat:
        interaction_parts.append(
```

```

        (X[rc] * X[cc]).rename(f'{rc}:{cc}')
    )
X_interact = pd.concat([X] + interaction_parts, axis=1)

print(f"Additive features: {X.shape[1]} columns")
print(f"Interaction features: {X_interact.shape[1]} columns "
      f"({X_interact.shape[1] - X.shape[1]} interaction terms)")

```

Additive features: 22 columns
Interaction features: 107 columns (+85 interaction terms)

Train/Test Split and Statsmodels Fit

Before fitting our models, we partition the dataset into an 80% training set and a 20% hold-out test set. Because our target variable is heavily skewed towards survival, we use a stratified split (`stratify=y`). This critical step ensures that both the training and testing sets preserve the exact 86.6% / 13.4% class proportions found in our full dataset, preventing the minority class from being underrepresented during training.

Using the training data, we then fit two logistic regression models using `statsmodels.Logit`. We choose this specific library for its comprehensive statistical output, which provides the p-values and confidence intervals required for rigorous inference:

The Additive Model: Captures the independent main effects of Region and Class on unit survival.

The Interaction Model: Incorporates our expanded feature matrix to test whether these two predictors interact, specifically, whether the survival probability of a certain commercial class fundamentally changes depending on its geographic location.

```

In [23]: X_train_int, X_test_int, y_train, y_test = train_test_split(
        X_interact, y, test_size=0.2, random_state=42, stratify=y
    )

additive_cols = X.columns.tolist()
X_train = X_train_int[additive_cols]
X_test = X_test_int[additive_cols]

X_train_sm = sm.add_constant(X_train)
X_test_sm = sm.add_constant(X_test)
sm_result = sm.Logit(y_train, X_train_sm).fit()

X_train_int_sm = sm.add_constant(X_train_int)
X_test_int_sm = sm.add_constant(X_test_int)
with warnings.catch_warnings():
    warnings.simplefilter('ignore')
    sm_interact = (sm.Logit(y_train, X_train_int_sm)
                  .fit(method='bfgs', disp=0, maxiter=500))

```

Optimization terminated successfully.
Current function value: 0.379660
Iterations 7

```

In [24]: sm_result.summary()

```

Dep. Variable:	Survived	No. Observations:	5369
Model:	Logit	Df Residuals:	5346
Method:	MLE	Df Model:	22
Date:	Fri, 27 Mar 2026	Pseudo R-squ.:	0.03837
Time:	10:40:25	Log-Likelihood:	-2038.4
converged:	True	LL-Null:	-2119.7
Covariance Type:	nonrobust	LLR p-value:	1.900e-23

	coef	std err	z	P> z	[0.025	0.975]
const	2.0190	0.135	14.962	0.000	1.755	2.284
Region_EH10	0.6813	0.246	2.775	0.006	0.200	1.163
Region_EH11	-0.2194	0.189	-1.164	0.245	-0.589	0.150
Region_EH12	-0.4572	0.201	-2.275	0.023	-0.851	-0.063
Region_EH13	0.5415	0.615	0.880	0.379	-0.664	1.747
Region_EH14	-0.6616	0.232	-2.853	0.004	-1.116	-0.207
Region_EH15	0.4007	0.293	1.369	0.171	-0.173	0.975
Region_EH16	0.2636	0.332	0.795	0.427	-0.386	0.913
Region_EH17	1.1058	0.736	1.502	0.133	-0.337	2.548
Region_EH2	0.0031	0.203	0.015	0.988	-0.394	0.400
Region_EH3	-0.2058	0.169	-1.215	0.224	-0.538	0.126
Region_EH30	-0.9915	0.394	-2.519	0.012	-1.763	-0.220
Region_EH4	-0.1280	0.215	-0.594	0.552	-0.550	0.294
Region_EH5	-1.3236	0.251	-5.280	0.000	-1.815	-0.832
Region_EH6	0.2667	0.196	1.359	0.174	-0.118	0.651
Region_EH7	0.5794	0.236	2.452	0.014	0.116	1.043
Region_EH8	-0.9270	0.173	-5.369	0.000	-1.265	-0.589
Region_EH9	0.3482	0.264	1.318	0.187	-0.170	0.866
Class_1986_CLASS 11	-0.5714	0.394	-1.451	0.147	-1.343	0.200
Class_1986_CLASS 2	-0.1079	0.146	-0.741	0.459	-0.393	0.178
Class_1986_CLASS 3	0.3707	0.219	1.694	0.090	-0.058	0.800
Class_1986_CLASS 4	-0.0172	0.154	-0.112	0.911	-0.318	0.284
Class_1986_SUI GENERIS	-0.2820	0.118	-2.385	0.017	-0.514	-0.050

Interpreting the Additive Model Summary

The statsmodels output provides a wealth of information about our model's performance and the impact of individual features. We can break this down into three key areas:

Overall Model Fit: The summary reports a McFadden's pseudo R^2 of 0.038. While the overall model is highly statistically significant (LLR p-value is effectively zero), this modest R^2 value confirms that

Region and Class alone explain only a limited share of the variation in commercial survival. This aligns with our expectations: important unmeasured variables such as rent, proximity to transport, area-level deprivation (e.g. the Scottish Index of Multiple Deprivation), and individual management decisions likely play substantial roles.

Coefficient Interpretation: Each coefficient (coef) represents the change in the log-odds of survival for that specific category relative to the dropped reference baseline. To make these values more intuitive, they can be exponentiated to yield the odds ratio.

Statistical Significance: We look at the $P > |z|$ column and the 95% confidence intervals ([0.025, 0.975]). Coefficients with a p-value < 0.05 and confidence intervals that completely exclude zero are deemed statistically significant predictors of long-term survival.

Odds Ratio Table

While log-odds coefficients are useful for model interpretation, **odds ratios** are more intuitive for decision-makers. An OR of 2.0 means the odds of survival are twice as high for that category relative to the reference; an OR of 0.5 means the odds are halved. We exponentiate each coefficient and its 95% confidence interval, sorting by OR to highlight the strongest predictors.

```
In [25]: # Clip coefficients to avoid overflow when exponentiating very large values
clipped_params = sm_result.params.clip(-20, 20)
clipped_ci      = sm_result.conf_int().clip(-20, 20)

odds_ratio_df = pd.DataFrame({
    'Odds Ratio': np.exp(clipped_params),
    'OR Lower 95%': np.exp(clipped_ci[0]),
    'OR Upper 95%': np.exp(clipped_ci[1]),
    'p-value': sm_result.pvalues,
}).drop('const')

# Flag statistical significance at the 5% level
odds_ratio_df['Significant'] = (
    odds_ratio_df['p-value']
    .lt(0.05)
    .map({True: 'Yes', False: 'No'})
)

odds_ratio_df.index = (
    odds_ratio_df.index
    .str.replace('Region_', 'Region: ', regex=False)
    .str.replace('Class_1986_', 'Class: ', regex=False)
)

odds_ratio_df = odds_ratio_df.sort_values('Odds Ratio', ascending=False)

odds_ratio_display = (odds_ratio_df
    .round({'Odds Ratio': 3, 'OR Lower 95%': 3,
           'OR Upper 95%': 3}).copy())

odds_ratio_display['p-value'] = odds_ratio_df['p-value'].map('{:.4f}'.format)
display(odds_ratio_display)
```

	Odds Ratio	OR Lower 95%	OR Upper 95%	p-value	Significant
Region: EH17	3.022	0.714	12.788	0.1330	No
Region: EH10	1.977	1.222	3.198	0.0055	Yes
Region: EH7	1.785	1.123	2.836	0.0142	Yes
Region: EH13	1.719	0.515	5.737	0.3786	No
Region: EH15	1.493	0.841	2.650	0.1711	No
Class: CLASS 3	1.449	0.943	2.225	0.0903	No
Region: EH9	1.417	0.844	2.377	0.1875	No
Region: EH6	1.306	0.889	1.918	0.1741	No
Region: EH16	1.302	0.680	2.493	0.4265	No
Region: EH2	1.003	0.674	1.492	0.9879	No
Class: CLASS 4	0.983	0.727	1.328	0.9109	No
Class: CLASS 2	0.898	0.675	1.194	0.4589	No
Region: EH4	0.880	0.577	1.342	0.5522	No
Region: EH3	0.814	0.584	1.134	0.2243	No
Region: EH11	0.803	0.555	1.162	0.2446	No
Class: SUI GENERIS	0.754	0.598	0.951	0.0171	Yes
Region: EH12	0.633	0.427	0.939	0.0229	Yes
Class: CLASS 11	0.565	0.261	1.222	0.1467	No
Region: EH14	0.516	0.328	0.813	0.0043	Yes
Region: EH8	0.396	0.282	0.555	0.0000	Yes
Region: EH30	0.371	0.172	0.803	0.0118	Yes
Region: EH5	0.266	0.163	0.435	0.0000	Yes

Interpretation: Odds Ratios

The Odds Ratio (OR) translates the model's coefficients into multiplicative effects on survival. An **OR > 1.0** indicates increased odds of survival relative to the baseline (**EH1** for Region and **Class 1** for Class), while an **OR < 1.0** indicates decreased odds.

- **Top Survival Predictors:** **EH10** (OR = 1.977) and **EH7** (OR = 1.785) are the strongest significant positive drivers. Units in EH10 are nearly twice as likely to survive as those in the baseline district.
- **Highest Risk Locations:** **EH5** (OR = 0.266) and **EH8** (OR = 0.396) show the most severe reduction in survival odds. A business in EH5 has roughly **73% lower odds** of surviving compared to the baseline. **EH30**, **EH14**, and **EH12** also exhibit significantly lower odds.
- **Commercial Classification:** **Sui Generis** units are the only class with a statistically significant effect ($p = 0.0171$), showing a **25% reduction** in survival odds (OR = 0.754). While **Class 3** (Food and Drink) trends positive (OR = 1.449), but does not meet the 5% significance threshold ($p = 0.0903$).

Summary: Geographic location (Region) offers a much wider range of significant multiplicative effects than commercial classification, reinforcing the conclusion that where a business is located is a more potent predictor of longevity than what it sells.

Balanced Sklearn Pipelines (Additive and Interaction)

The `statsmodels` models are valuable for inference but do not address class imbalance: with an 86.6% baseline survival rate, an unweighted model tends to predict "survived" for nearly every unit. We wrap `LogisticRegression(class_weight='balanced')` in an `sklearn.Pipeline` with `StandardScaler` to up-weight the minority class and guarantee consistent feature scaling across train and test sets.

```
In [26]: # Additive pipeline
pipe = make_pipeline(
    StandardScaler(),
    LogisticRegression(class_weight='balanced', max_iter=1000, random_state=42)
)
pipe.fit(X_train, y_train)
y_pred = pipe.predict(X_test)
y_pred_prob = pipe.predict_proba(X_test)[: , 1]

# Interaction pipeline
pipe_interact = make_pipeline(
    StandardScaler(),
    LogisticRegression(class_weight='balanced', max_iter=1000, random_state=42)
)
pipe_interact.fit(X_train_int, y_train)
y_pred_int = pipe_interact.predict(X_test_int)
y_pred_prob_int = pipe_interact.predict_proba(X_test_int)[: , 1]
```

Model Evaluation on the Test Set

We evaluate both the additive and interaction pipelines on the held-out test set. Because `class_weight='balanced'` penalises errors on the minority class more heavily, we expect improved recall for non-surviving units, even if overall accuracy drops slightly. Comparing the two pipelines reveals whether interaction terms translate into better predictive performance.

```
In [27]: print("=" * 60)
print("ADDITIVE MODEL (Region + Class)")
print("=" * 60)
print(f"Accuracy: {accuracy_score(y_test, y_pred):.1%}")
print(classification_report(y_test, y_pred,
                           target_names=['Did Not Survive', 'Survived']))

print("=" * 60)
print("INTERACTION MODEL (Region × Class)")
print("=" * 60)
print(f"Accuracy: {accuracy_score(y_test, y_pred_int):.1%}")
print(classification_report(y_test, y_pred_int,
                           target_names=['Did Not Survive', 'Survived']))
```

ADDITIVE MODEL (Region + Class)

Accuracy: 59.2%

	precision	recall	f1-score	support
Did Not Survive	0.19	0.62	0.29	180
Survived	0.91	0.59	0.71	1163
accuracy			0.59	1343
macro avg	0.55	0.60	0.50	1343
weighted avg	0.81	0.59	0.66	1343

INTERACTION MODEL (Region × Class)

Accuracy: 61.6%

	precision	recall	f1-score	support
Did Not Survive	0.20	0.61	0.30	180
Survived	0.91	0.62	0.74	1163
accuracy			0.62	1343
macro avg	0.55	0.61	0.52	1343
weighted avg	0.82	0.62	0.68	1343

Understanding the Classification Report

To evaluate the model's effectiveness, we examine precision (the proportion of predicted failures that actually failed), recall (the proportion of actual failures that the model detected), and the F1-score (their harmonic mean).

Given the significant class imbalance in our test set, where only 180 out of 1,343 units did not survive overall accuracy is a misleading metric. A naive model that simply predicted "Survived" for every unit would achieve 86.6% accuracy but have 0% recall for the units we are most interested in identifying.

By using balanced class weighting, the additive model achieved a recall of 0.62 for the "Did Not Survive" class. This means the model successfully identified 62% of all units that actually failed. While this trades off overall accuracy (dropping it to 59%) and results in lower precision (0.19), it is the correct strategy for this research. In commercial survival analysis, it is far more valuable to flag a high-risk unit (even if it eventually survives) than to miss it entirely.

Coefficient Plot

This forest plot displays each estimated log-odds coefficient with its 95% confidence interval, enabling rapid visual comparison of all predictors. Coefficients whose intervals exclude the dashed zero line are statistically significant at the 5% level. We exclude the intercept from the plot since its scale differs and it does not correspond to a specific region or use-class.

```
In [28]: label_map = {
    'Region_EH1': 'EH1 - City Centre',
    'Region_EH2': 'EH2 - City Centre',
    'Region_EH3': 'EH3 - Inner City',
    'Region_EH4': 'EH4 - North West',
    'Region_EH5': 'EH5 - North (Coastal)',
    'Region_EH6': 'EH6 - North East',
```

```

'Region_EH7': 'EH7 - Inner East',
'Region_EH8': 'EH8 - Inner East',
'Region_EH9': 'EH9 - Inner South',
'Region_EH10': 'EH10 - South',
'Region_EH11': 'EH11 - Inner West',
'Region_EH12': 'EH12 - West',
'Region_EH13': 'EH13 - South West',
'Region_EH14': 'EH14 - Outer South West',
'Region_EH15': 'EH15 - East (Coastal)',
'Region_EH16': 'EH16 - South East',
'Region_EH17': 'EH17 - Outer South',
'Region_EH30': 'EH30 - West Outskirts',

'Class_1986_CLASS 2': 'Class 2 - Financial & Professional Services',
'Class_1986_CLASS 3': 'Class 3 - Food & Drink',
'Class_1986_CLASS 4': 'Class 4 - Business',
'Class_1986_CLASS 11': 'Class 11 - Assembly & Leisure',
'Class_1986_SUI GENERIS': 'Sui Generis - Unique/Other',
}

# Process coefficients
coef_df = (
    pd.DataFrame({
        'Coefficient': sm_result.params,
        'Lower CI': sm_result.conf_int()[0],
        'Upper CI': sm_result.conf_int()[1],
    })
    .drop('const')
    .sort_values('Coefficient')
)

# Filter out extreme outliers with massive confidence intervals
ci_width = coef_df['Upper CI'] - coef_df['Lower CI']
coef_df = coef_df[ci_width < 50]

coef_df.index = coef_df.index.map(lambda x: label_map.get(x, x))

cb = sns.color_palette("colorblind")
colours = [cb[0] if 'EH' in lab else cb[3] for lab in coef_df.index]

coeffs = coef_df['Coefficient'].values
err_lower = coeffs - coef_df['Lower CI'].values
err_upper = coef_df['Upper CI'].values - coeffs
x_pos = list(range(len(coef_df)))

# Plotting
fig, ax = plt.subplots(figsize=(20, 10))

for i, (coeff, el, eu, col) in enumerate(
    zip(coeffs, err_lower, err_upper, colours)):
    ax.errorbar(x=i, y=coeff, yerr=[[el], [eu]],
                fmt='o', color=col, ecolor=col,
                elinewidth=2, capsize=4, capthick=2, markersize=7)

ax.axhline(y=0, color=cb[6], linestyle='--', linewidth=2.0)

ax.set_xticks(x_pos)
ax.set_xticklabels(coef_df.index, rotation=50, ha='right', fontsize=18)
ax.set_ylabel('Estimated Coefficient (log-odds of survival)', fontsize=20)
ax.set_xlabel('Predictor Variable', fontsize=24)
ax.set_title('Logistic Regression Coefficients with 95% Confidence Intervals',
             fontsize=24, pad=12)
ax.tick_params(axis='y', labelsiz=20)

legend_handles = [
    Line2D([0], [0], marker='o', color=cb[0], linestyle='None',


```

```

        markersize=8, label='Region'),
        Line2D([0], [0], marker='o', color=cb[3], linestyle='None',
                markersize=8, label='Commercial class'),
        Line2D([0], [0], color=cb[6], linestyle='--',
                linewidth=2.0, label='Zero effect'),
    ]
ax.legend(handles=legend_handles, loc='upper left', fontsize=18)

fig.tight_layout()
fig.savefig('coefficient_plot.png', dpi=200, bbox_inches='tight')
plt.show()

```



Interpretation: This figure mainly confirms inferential results already reported above: several Region coefficients are clearly away from zero, while most Class effects are smaller with wider overlap around zero. The key value of the plot is visual, it shows at a glance that Region effects are both stronger and more consistently significant than Class effects.

Region vs. Commercial Class — Which Matters More?

To answer our core research question, we compare nested logistic regression models using the **Likelihood Ratio Test (LRT)**. Dropping Region from the full additive model should produce a larger LRT statistic than dropping Class if location is the stronger predictor. We also test whether adding Region × Class interaction terms significantly improves fit, which would indicate that the two predictors do not operate independently.

```

In [29]: # Identify column groups
region_cols = [c for c in X_train.columns if c.startswith('Region_')]
class_cols = [c for c in X_train.columns if c.startswith('Class_1986_')]

# Fit restricted models on training data (with statsmodels constant)
with warnings.catch_warnings():
    warnings.simplefilter('ignore')
    region_only = sm.Logit(y_train, sm.add_constant(
        X_train[region_cols])).fit(dis=0, maxiter=500)
    class_only = sm.Logit(y_train, sm.add_constant(
        X_train[class_cols])).fit(dis=0, maxiter=500)

full_ll = sm_result.llf
reg_ll = region_only.llf
cls_ll = class_only.llf

```

```

interact_ll = sm_interact.llf

# LRT: additive vs restricted
lr_drop_class = -2 * (reg_ll - full_ll)
lr_drop_region = -2 * (cls_ll - full_ll)
df_class = len(class_cols)
df_region = len(region_cols)
p_drop_class = chi2.sf(lr_drop_class, df_class)
p_drop_region = chi2.sf(lr_drop_region, df_region)

# LRT: interaction vs additive
interaction_cols = [c for c in X_train_int.columns if ':' in c]
lr_interaction = -2 * (full_ll - interact_ll)
df_interaction = len(interaction_cols)
p_interaction = chi2.sf(lr_interaction, df_interaction)

# Comparison table
comparison = pd.DataFrame({
    'LRT Statistic': [None, lr_drop_class, lr_drop_region, lr_interaction],
    'LRT p-value': [None, p_drop_class, p_drop_region, p_interaction],
}, index=[
    'Full (Region + Class)',
    'Region-only (drop Class)',
    'Class-only (drop Region)',
    'Interaction (Region × Class)',
])

# Format
comparison_display = comparison.copy()
comparison_display['LRT Statistic'] = comparison_display['LRT Statistic'].apply(
    lambda x: f'{x:.2f}' if pd.notna(x) else '-')
comparison_display['LRT p-value'] = comparison_display['LRT p-value'].apply(
    lambda x: f'{x:.2e}' if pd.notna(x) else '-')

print("Model Comparison – Region vs. Commercial Class vs. Interaction\n")
display(comparison_display)

```

Model Comparison – Region vs. Commercial Class vs. Interaction

	LRT Statistic	LRT p-value
Full (Region + Class)	—	—
Region-only (drop Class)	11.61	4.05e-02
Class-only (drop Region)	150.82	1.69e-23
Interaction (Region × Class)	102.69	9.30e-02

Interpretation: The LRT decisively answers our core research question. Dropping Region from the full model produces a far larger LRT statistic (**150.82**, $p = 1.69 \times 10^{-23}$) than dropping Class (**11.61**, $p = 0.0405$), confirming that **geographic location is the substantially stronger predictor** of long-term survival.

Adding Region × Class interaction terms does *not* significantly improve model fit (LRT = 102.69, $p = 0.0930$), indicating the two predictors operate **independently**. An additive model is sufficient: certain regions are harder for *all* business types, and certain classes are more vulnerable *everywhere*.

ROC Curve and AUC


The ROC curve evaluates model discrimination across all classification thresholds rather than a single arbitrary cutoff, which is especially valuable given our class imbalance. An AUC significantly above 0.5 confirms the model has learned meaningful patterns.

```
In [30]: # ROC curves – Additive vs. Interaction
fpr_add, tpr_add, _ = roc_curve(y_test, y_pred_prob)
roc_auc_add = auc(fpr_add, tpr_add)

fpr_int, tpr_int, _ = roc_curve(y_test, y_pred_prob_int)
roc_auc_int = auc(fpr_int, tpr_int)

cb = sns.color_palette("colorblind")
fig, ax = plt.subplots(figsize=(7, 6))
ax.plot(fpr_add, tpr_add, color=cb[0], lw=2,
        label=f'Additive model (AUC = {roc_auc_add:.3f})')
ax.plot(fpr_int, tpr_int, color=cb[3], lw=2, linestyle='-.-',
        label=f'Interaction model (AUC = {roc_auc_int:.3f})')
ax.plot([0, 1], [0, 1], color=cb[6], lw=1.5, linestyle='--',
        label='Random classifier (AUC = 0.500)')
ax.set_xlim([0.0, 1.0])
ax.set_ylim([0.0, 1.05])
ax.set_xlabel('False Positive Rate (1 – Specificity)', fontsize=18)
ax.set_ylabel('True Positive Rate (Sensitivity)', fontsize=18)
ax.set_title('ROC Curve – Additive vs. Interaction Model (Test Set)', fontsize=20)
ax.legend(loc='lower right', fontsize=16)
ax.tick_params(axis='both', labelsize=18)
fig.tight_layout()
fig.savefig('roc_curve.png', bbox_inches='tight')
plt.show()
```

ROC Curve — Additive vs. Interaction Model (Test Set)



Interpretation: Both models perform above the random-classifier diagonal ($AUC > 0.5$), confirming they have learned meaningful patterns from Region and Class alone. However, the modest test AUC of **0.655** indicates limited discriminative power. The model correctly ranks a randomly chosen surviving unit above a non-surviving one only about 65.5% of the time. The interaction model shows no meaningful AUC improvement over the additive model, reinforcing the LRT finding that interaction terms are unnecessary.

This modest performance is expected: Region and Class capture broad geographic and sectoral effects, but unmeasured confounders — such as rent levels, proximity to transport, area-level deprivation, and individual management decisions — likely account for much of the remaining variation in survival outcomes.

I used AI (Gemini) to help with clarity within my writing